!mpactmakers

A/B = A Bad Idea?

Improved Insights with Design & Analysis of Experiments in R

David Moxley

11 March 2020

Intro

David Moxley
Senior Consultant
Impact Makers

Founded in **2006**

$20+ million revenue

80 employees

Certified B Corp

Motivation

- Lack of emphasis on experimentation in the Richmond market
- Prevalence in a variety of industries and business units
- Popularity of A/B design
- Ignored subtleties of A/B tests

Agenda

- Why do we experiment & why analyze in R?
- What's an A/B experiment?
- Case Studies
- Alternative Design of Experiments
- Analyzing in R
- Developing a broader "Test and Learn" culture

01 | Why R?

Tools

The Marketplace

General-Purpose

Purpose-Built

Statistical Languages	DoE Analysis Tool	Automated General DoE	Subject-Specific
R	jmp Statistical Discovery.™ From SAS.	APT A Mastercard Company	Optimizely

Technical Demand of User

Why analyze experiments in R?

As a practitioner…

- Continuity with other workflows
- Transparency
- Powerful visualization packages
- Instructional value*

As a business leader…

- Extensible framework
- Transparency
- Learning curve for analytical teams
- Cost

02 | Why do we experiment?

Why do we experiment?

- "You break you buy it"

Why do we experiment?

- ~~"You break you buy it"~~
- Break it before you buy it
- Foster a learning culture
- Tease out causality
- Provide direction to the business

The Goal

- Minimize the impact on the business*
- Offer a simple design
- Eliminate systematic error
- Understand the range of validity
- Offer a precise estimate
- Convey uncertainty
- Iterate!

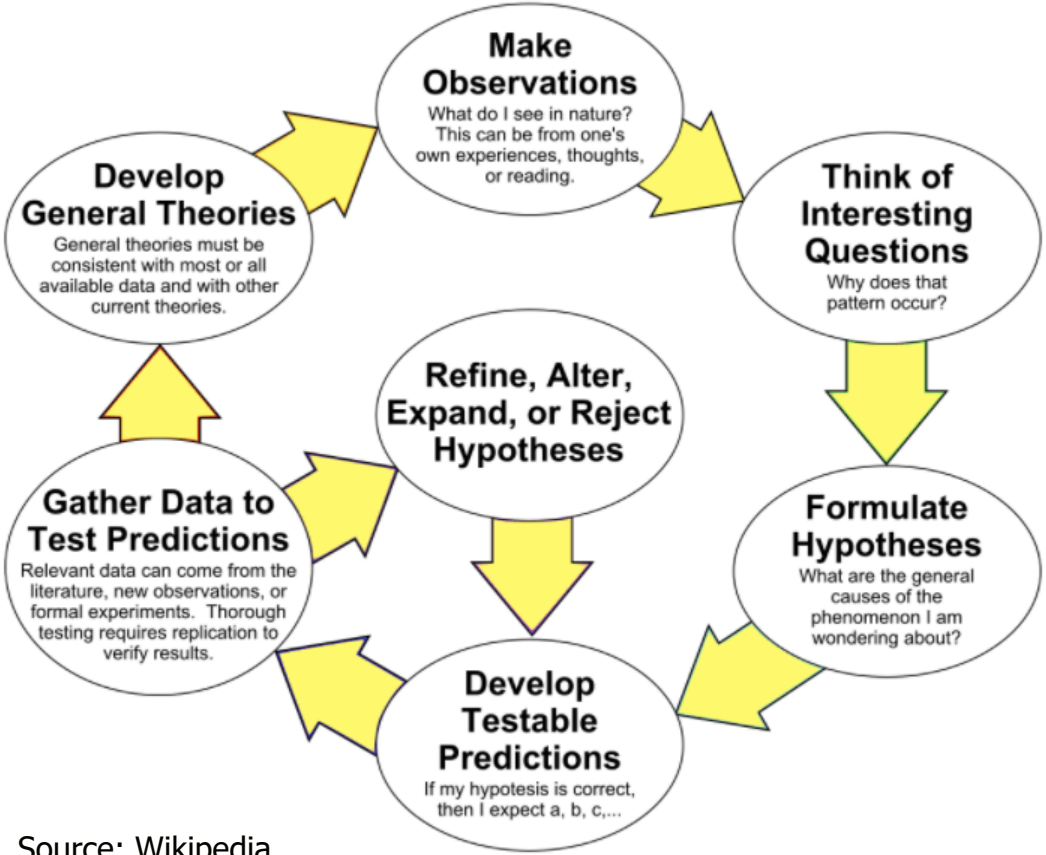

Source: Wikipedia

Types of Experiments

Types	Example	Strengths	Potential Issues
Laboratory	• Survey Research	• High internal validity • Ease of replicability	• Lack of Realism • Poor Generalizability
Field	• Tele-marketing	• Strong internal validity • Very Realistic	• Generalizability (?) • Potential selection bias
Natural	• TV Ad test	• Highly Realistic • Generalizability (?)	• Causal Inference (?) • Data Collection

What's an A/B Experiment?

- Randomized Controlled Trial
    - Split-run testing
- Random assignment to two or more variants, A, B,…*n*
- Widely used for testing Machine Learning models
- Popular in digital space, UX research, etc.

03 | Naïve A/B Designs

Case Studies

Case Study: Price Elasticity Test

- B2B retailer tested price changes for a set of SKUs
    - Cell 1: 5% increase, free shipping
    - Cell 2: 10% increase, free shipping
- Salespeople/Accounts pseudo-randomly assigned to cells
    - Cells were balanced for total sales
- Total Sales for the Cells were measured
- Analysis via Difference in Differences
- Test continually monitored until predetermined alpha of .1 was reached

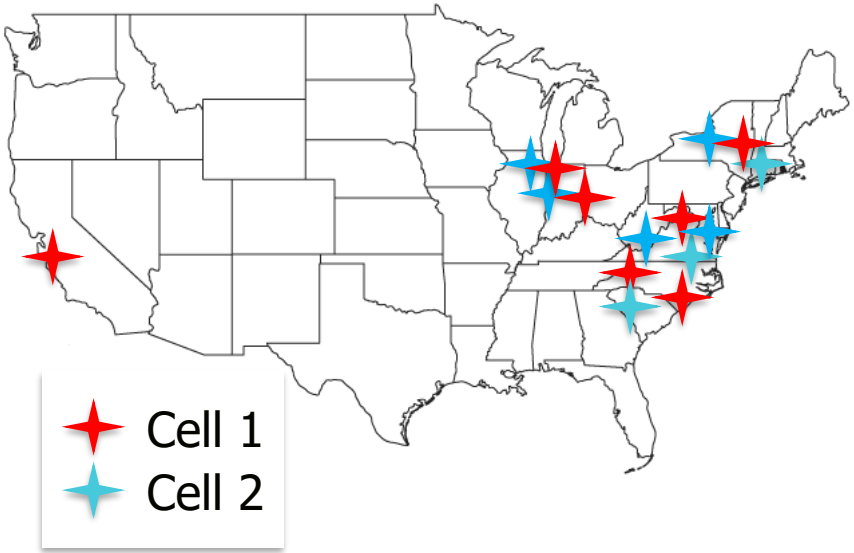

Case Study: Potential Drawbacks

- No treatment for confounding factors
- Sample bias invalidated Difference in Differences analysis
- Entangle effects through overly simple design
    - No understanding of interactions
- "Peeking" violated an underlying assumption of statistical inference

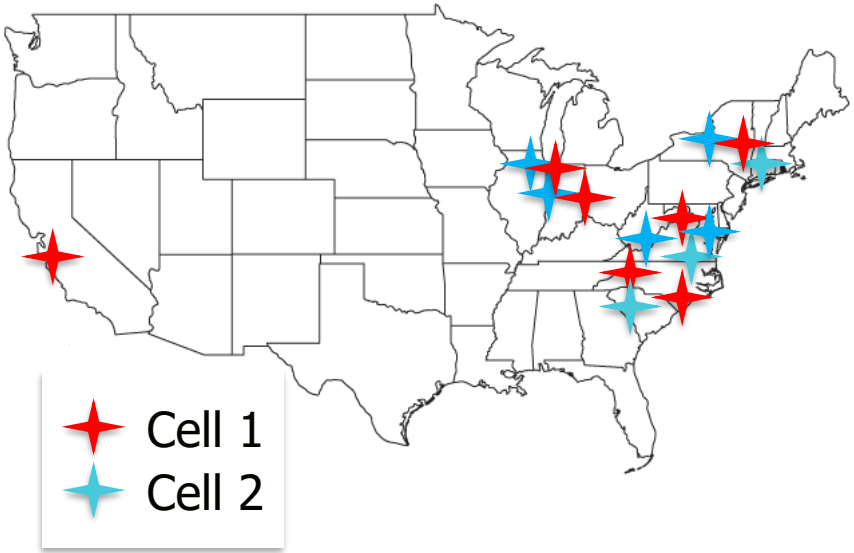

What is “Peeking”?

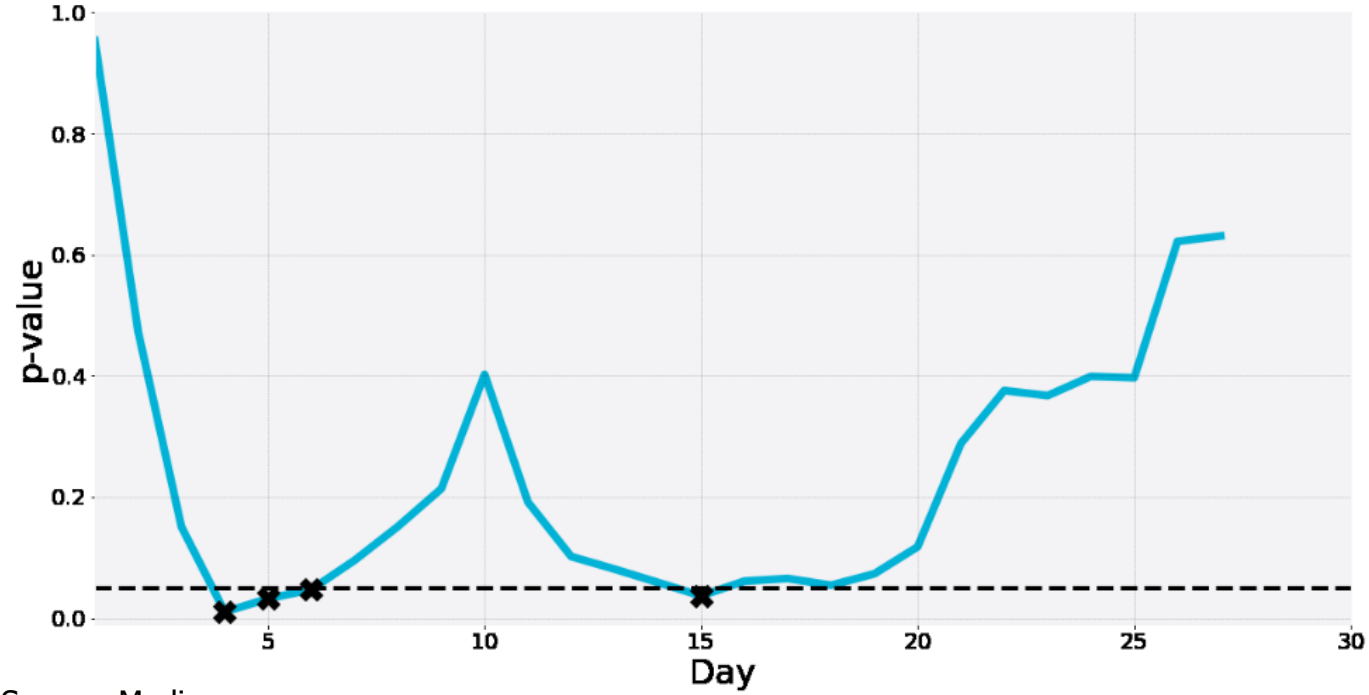

Source: Medium.com

The Real Issue

"To consult the statistician[/data scientist/consultant] after an experiment is finished is often merely to ask him to conduct a post-mortem examination. He can perhaps say what the experiment died of."

-R.A. Fisher, 1938

04 | Design of Experiments

Overview

Keys to a Good Design

1. Replication
2. Randomization
3. Control

Types of Designs

- Comparative studies
- Single Factor
- Blocking Designs
  - Randomized Complete Block
  - Balanced Incomplete Block
- Factorial
- Fractional Factorial
- Response Surface Designs

Types of Designs

- Comparative studies
- Single Factor
- Blocking Designs
  - Randomized Complete Block
  - Balanced Incomplete Block
- Factorial
- Fractional Factorial
- Response Surface Designs

05 | Redesign Our Case Study

Case Study: Two-Factor Factorial in a Randomized Complete Block

Sales Person	Price % Increase	Shipping Fee
1	5	Yes
1	10	No
1	10	Yes
1	5	No
…	…	
5	5	Yes
5	10	No
5	10	Yes
5	5	No

$$y_{ijk} = \mu + \tau_i + \alpha_j + (\tau\alpha)_{ij} + \beta_k + \epsilon \begin{cases} i = 1 \; to \; 2 \\ j = 1 \; to \; 2 \\ k = 1 \; to \; 5 \end{cases}$$

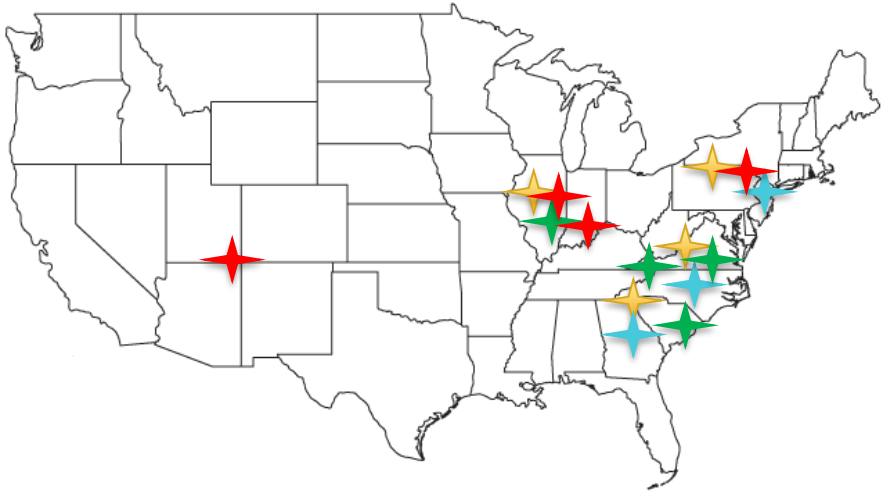

Case Study: Sequential Sampling Procedure

- Frequentist hypothesis testing assumes **fixed sample**
  - Central Limit Theorem
- Business pressure, ethical considerations, user negligence can lead to a desire to monitor, "peek" at test results
- Larger conversation intersects with Theory of Optimal Stopping
  - The "Secretary Problem"

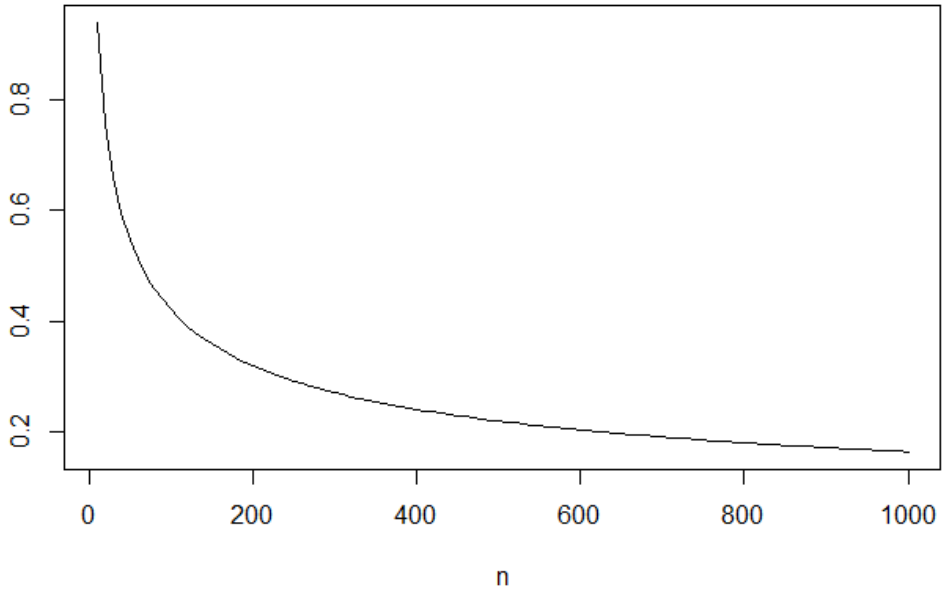

Case Study: Sequential Sampling Procedure

- Wald's Sequential Probability Ratio Test (SPRT)
- Optimizely's Mixture Sequential Probability Test (mSPRT)
- Multi-Armed Bandit
- Bayesian Methods
- Evan Miller's Sequential Procedure with Stopping Metric
    1. Choose a target sample size, n, at the outset
    2. Assign to treatments with equal probabilities
    3. Track incoming successes for each treatment cell, A, B,..*i*
    4. If A – B = $2\sqrt{n}$, declare A the winner, B-A = $2\sqrt{n}$, B is the winner
    5. If T + C = N, there is no winner, fail to reject null hypothesis

06 | Analysis Procedures

Overview

Stats Refresher

- Degrees of Freedom
- T-test
  - One-Sample: $\frac{\bar{X}-u}{\frac{sd}{\sqrt{n}}}$
  - Unpaired: $\frac{\bar{X}_1-\bar{X}_2}{sd_p\sqrt{\frac{1}{n_1}-\frac{1}{n_2}}}$
  - Paired: $\frac{\bar{X}_d}{\frac{sd_d}{\sqrt{n}}}$
- Sum of Squares: $\sum(x_i-\bar{x})^2$
- F-test
  - $\frac{MSE_{Larger\ Sample}}{MSE_{Smaller\ Sample}}$
- p-value: Type I error
- Power: Type II error

Stats Refresher

- ANOVA

$$y_{ijk} = \mu + \tau_i + \alpha_j + \beta_k + \epsilon \begin{cases} i = 1 \text{ to } i \\ j = 1 \text{ to } j \\ k = 1 \text{ to } k \end{cases}$$

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
A	1	1116	1116	387.430	< 2e-16	***
B	1	9214	9214	3197.928	< 2e-16	***
C	1	751	751	260.575	9.88e-15	***
D	1	5	5	1.833	0.188	
E	1	2	2	0.531	0.473	
A:B	1	504	504	174.935	8.68e-13	***
Residuals	25	72	3			

General Procedure

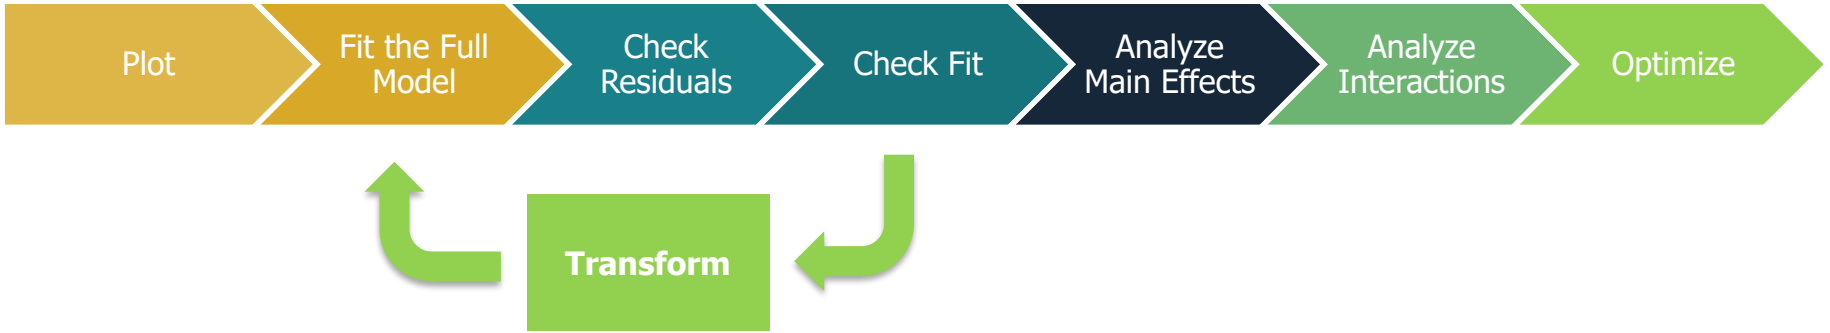

07 | Analyzing in R

Case Studies

R Script

08 | A/B[ad] Idea?

Conclusion

Are A/B Designs a good idea?

Pros	Cons
Simple Design	Entangled Effects
Statistical efficiency with minimal sample	Susceptible to systematic bias
Well suited for digital studies	Potential issues with “peeking”

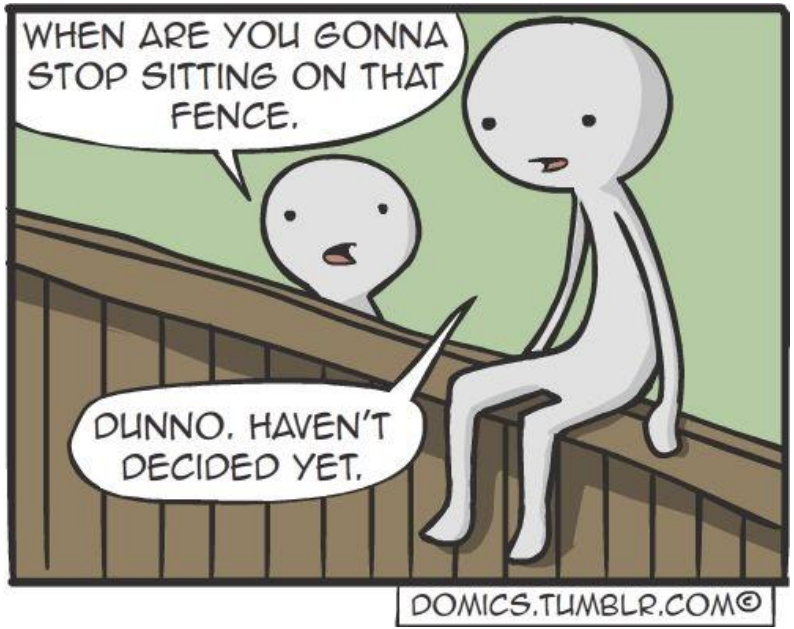

Parting Advice

- Collaborate with the business
- Keep Design and Analysis simple
- Remember statistical vs business significance
- Experimentation is mean to be iterative

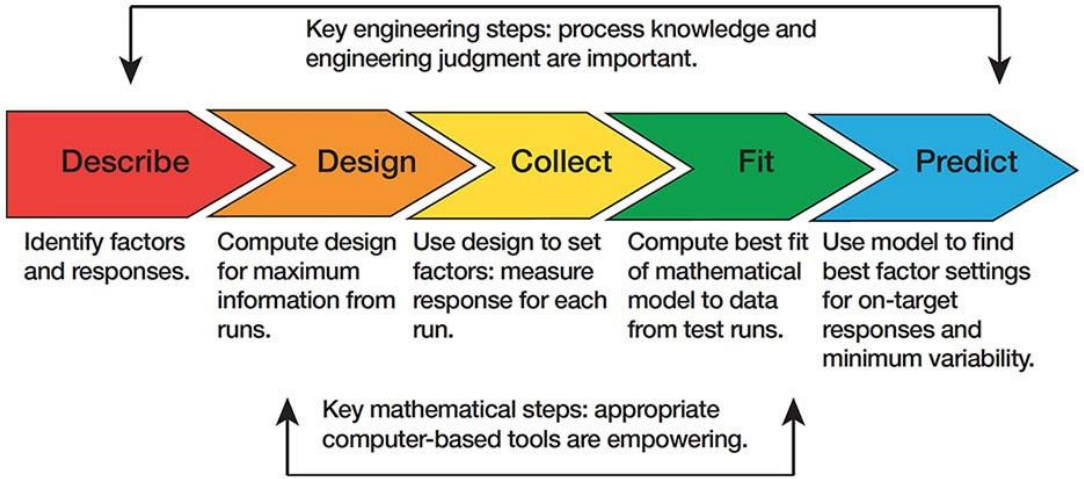

Source: JMP.com

Thank you!

- Connect: dmoxley@impactmakers.com
- Questions: dmoxley@impactmakers.com
- Data/Slides/R Code: github.com/davidrmoxley/DoE
- Continue the Conversation: impactmakers.lpages.co/advanced-analytics/